

Static Analysis for TOCTTU bugs

I. CASE STUDY

A. Case 1 (*Suspicious*)

- Global Variable: `acpi_gbl_next_history_index`
- File: `drivers/acpi/acpica/dbhistory.c`
- Function: `acpi_db_add_to_history`
- Use Location 1: Line 90 Use as index
- Use Location 2: Line 94 Use as index
- Use Location 3: Line 99 Use as index
- Use Location 4: Line 100 Use as index
- Use Location 5: Line 104 Use as index
- Use Location 6: Line 108 Use as index
- Use Location 7: Line 117 Use as check
- Use Location 8: Line 125 Use as check
- Define Location 1: Line 124 Reset
- Define Location 2: Line 126 Increment
- Analysis;
 - Use/Def Analysis:
 - Use as array index and serve condition check at line 90; May cause undefined behavior when corrupted
 - Use as array index and serve string length calculation at line 94; May pollute the string length and further memory overflow
 - Use as array index and serve memory copy at line 108; May lead to arbitrary memory write
 - Lock Analysis
 - Reverse call sequence: `acpi_db_add_to_history` (No lock) → `acpi_db_command_dispatch` (No lock) → `acpi_db_user_commands` (No lock) → `acpi_db_execute_thread` (No lock) → `acpi_os_execute` (Being started as a new thread; No lock)
 - Parallel Access Analysis:
 - Exploitability Analysis

B. Case 2 (*Suspicious*)

- Global Variable: `he_devs`
- File: `drivers/atm/he.c`
- Function: `he_init_one`
- Use Location 1: Line 395 Use as check

- Use Location 2: Line 396 Use as node on linked list
- Define Location 2: Line 397 define as the new head of linked list
- Analysis;
 - Use/Def Analysis:
 - Use as the current head of a linked list at Line 395
 - Update it to the new head the current linked list at Line 396
 - Lock Analysis
 - Reverse call sequence: `he_init_one` (No lock) → used as the probe handler of the "he" pci driver.
 - Parallel Access Analysis:
 - Exploitability Analysis: If two threads can reach Line 397 at the same time, then only one thread can update `he_devs` (aka, the head of the linked list). As a consequence, one newly allocated node is not linked in and can never be freed (memory leakage)?

C. Case 3 (*Unknown; Looks less interesting*)

- Global Variable: `iadev_count`
- File: `drivers/atm/iphase.c`

D. Case 4 (*Similar to Case 2*)

- Global Variable: `ia_boards`
- File: `drivers/atm/iphase.c`
- Function: `ia_init_one`

E. Case 5 (*Similar to Case 3*)

- Global Variable: `num_cards`
- File: `drivers/atm/nicstar.c`

F. Case 6 (*Not very interesting*)

- Global Variable: `fpga_upgrade`
- File: `drivers/atm/solos-pci.c`

G. Case 7 (*Similar to Case 2*)

- Global Variable: `eni_boards`
- File: `drivers/atm/eni.c`
- Function: `eni_init_one`

II. FALSE POSITIVE FILTERS FOR A PAIR OF USES

A. Reachability

The pair of uses are in at least one execution path (reachable in the CFG)

B. Lock

The pair of use are not covered by a set of lock/unlock

C. Define

At least one define can be affected by external input

D. Initialization

The pair of uses are not in an initialization function (Check if a function is placed into the .init.text section)

The define is not in an initialization function (Check if a function is placed into the .init.text section)

III. COUNTER EXAMPLES

A. Class 1: lock is added by a parent (still a dominator)

```
int device_attach(struct device *dev)
{
    int ret = 0;
    device_lock(dev);
    if (dev->driver) {
        if (klist_node_attached(&dev->p->knode_driver)) {
            ret = 1;
            goto out_unlock;
        }
        ret = device_bind_driver(dev);
        if (ret == 0)
            ret = 1;
        else {
            dev->driver = NULL;
            ret = 0;
        }
    } else {
        ret = bus_for_each_drv(dev->bus, NULL, dev, device_attach);
        pm_request_idle(dev);
    }
out_unlock:
    device_unlock(dev);
    return ret;
}
```

B. Class 2: The function is only called by init/exit function (in principle, same as init/exit function)

```
int xs_init(void)
{
    int err;
    struct task_struct *task;

    INIT_LIST_HEAD(&xs_state.reply_list);
    spin_lock_init(&xs_state.reply_lock);
    init_waitqueue_head(&xs_state.reply_waitq);
    xsduld_unregister_test(unsigned ioctl)
```

```
mutex_init(&xs_state.request_mutex);
mutex_init(&xs_state.response_mutex);
mutex_init(&xs_state.transaction_mutex);
init_rwsem(&xs_state.watch_mutex);
atomic_set(&xs_state.transaction_count, 0);
init_waitqueue_head(&xs_state.transaction_waitq);
```

```
/* Initialize the shared memory rings to
err = xsb_init_comms();
if (err)
    return err;
```

```
task = kthread_run(xenwatch_thread, NULL, "xenwatch");
if (IS_ERR(task))
    return PTR_ERR(task);
xenwatch_pid = task->pid;
```

```
task = kthread_run(xenbus_thread, NULL, "xenbus");
if (IS_ERR(task))
    return PTR_ERR(task);
```

```
return 0;
```

C. Class 3: Lock/unlock during the atomicity operation is legit

```
static void *read_reply(enum xsd_sockmsg_type *type, char *body)
{
    struct xs_stored_msg *msg;
    spin_lock(&xs_state.reply_lock);
    while (list_empty(&xs_state.reply_list))
        spin_unlock(&xs_state.reply_lock);
    /* XXX FIXME: Avoid synchronous wait */
    wait_event(xs_state.reply_waitq, !list_empty(&xs_state.reply_list));
    spin_lock(&xs_state.reply_lock);
    msg = list_entry(xs_state.reply_list.next, struct xs_stored_msg, list);
    list_del(&msg->list);
    spin_unlock(&xs_state.reply_lock);
    *type = msg->hdr.type;
    if (len)
        *len = msg->hdr.len;
    body = msg->u.reply.body;
    kfree(msg);
    return body;
}
```

D. Class 4: I have no idea here...

```
void xsduld_unregister_test(unsigned ioctl)
```

```
{  
    if ( ioctl == g_test_ioctl) {  
        g_test_ioctl = 0;  
        g_do_test = NULL;  
    }  
}
```

IV. NOTE

hfi1_user_sdma_process_request

does not check size after copy from user.